Enterprise-Level MCQs – Generative AI & LLM Foundations (Answers)

Q1. Enterprise AI Assistant Rollout

A. The model is performing deterministic rule-based execution only

B. The model is suffering from hallucination due to probabilistic token prediction

C. The tokenizer is incorrectly converting embeddings into vectors

D. The transformer architecture cannot process enterprise datasets

Answer = B

Q2. Transformer Architecture in Enterprise Search

A. Sequential memory storage using hidden states only

B. Recursive statistical loops

C. Self-attention mechanism enabling parallel context processing

D. Static keyword mapping tables

Answer = C

Q3. Tokenization Challenge in Healthcare AI

A. To permanently compress the database size

B. To convert raw text into processable numerical units for the model

C. To encrypt patient records before inference

D. To remove semantic relationships from language

Answer = B

Q4. Enterprise Inference Optimization

A. Increasing token length for all prompts

B. Using quantization and optimized inference hardware

C. Replacing transformers with spreadsheets

D. Expanding prompt size indefinitely

Answer = B

Q5. Generative AI Use Case Identification

A. Predictive analytics only

B. Rule-based automation only

C. Generative text synthesis using contextual enterprise data

D. Traditional database normalization

Answer = C

Q6. Context Window Limitation

A. Transformer models cannot process insurance documents

B. The context window limitation prevents retaining all tokens simultaneously

C. Embeddings delete historical information automatically

D. The tokenizer only supports numerical data

Answer = B

Q7. Enterprise LLM Deployment Strategy

A. Availability of social media integration

B. GPU cabinet color

C. Data privacy, compliance, and deployment control

D. Number of emojis supported

Answer = C

Q8. Attention Mechanism Scenario

A. Fixed rule engine

B. Self-attention relationships between tokens

C. Relational database joins

D. Data deduplication indexing

Answer = B

Q9. Hallucination Risk Management

A. Removing enterprise documents from prompts

B. Integrating Retrieval-Augmented Generation (RAG)

C. Increasing randomness temperature to maximum

D. Disabling tokenization

Answer = B

Q10. Enterprise GPU Infrastructure

A. GPUs contain built-in databases

B. GPUs are optimized for massive parallel mathematical operations

C. CPUs cannot run Python

D. GPUs eliminate tokenization requirements

Answer = B

Q11. Zero-Shot Prompting in HR Automation

A. Fine-tuning

B. Zero-shot prompting

C. Reinforcement learning

D. Vector indexing

Answer = B

Q12. Few-Shot Prompting Scenario

A. It permanently retrains the model

B. It reduces model parameter size

C. It guides the model toward desired response patterns

D. It converts the transformer into an RNN

Answer = C

Q13. Chain-of-Thought Reasoning

A. The prompt activates external APIs automatically

B. Chain-of-thought prompting encourages intermediate reasoning steps

C. The tokenizer becomes more efficient

D. The vector database automatically retrains the model

Answer = B

Q14. Prompt Optimization in Enterprise Analytics

A. Using ambiguous business terms intentionally

B. Providing clear instructions, constraints, and expected output format

C. Minimizing all contextual information

D. Avoiding structured prompts

Answer = B

Q15. Prompt Injection Risk

A. Semantic chunking

B. Prompt injection attack

C. Embedding normalization

D. Vector pruning

Answer = B

Q16. Structured Prompting for JSON Output

A. Using random conversational prompts

B. Defining strict formatting instructions and schema examples

C. Increasing token randomness

D. Removing output constraints entirely

Answer = B

Q17. Temperature Parameter Scenario

A. Responses become more deterministic

B. Responses become more creative and diverse

C. Tokenization stops functioning

D. Context windows increase automatically

Answer = B

Q18. Enterprise Prompt Evaluation

A. Database sharding

B. Prompt evaluation and optimization

C. Hardware virtualization

D. Token compression

Answer = B

Q19. Role-Based Prompting

A. Reducing GPU memory usage

B. Assigning contextual behavioral guidance to the model

C. Replacing embeddings

D. Encrypting transformer weights

Answer = B

Q20. Prompt Chaining in Enterprise Workflow

A. Prompt chaining

B. Data normalization

C. Token truncation

D. Recursive indexing

Answer = A

Q21. Semantic Search Implementation

A. They permanently store PDF files

B. They represent semantic meaning in numerical vector space

C. They replace transformer models entirely

D. They remove legal terminology from documents

Answer = B

Q22. Similarity Search Scenario

A. Bubble sort complexity

B. Cosine similarity

C. Binary tree depth

D. TCP packet latency

Answer = B

Q23. Vector Database Use Case

A. Traditional relational database only

B. Vector database optimized for nearest-neighbor search

C. Spreadsheet macros

D. FTP server indexing

Answer = B

Q24. Embedding Model Scenario

A. To perform semantic matching between related concepts

B. To reduce internet bandwidth usage only

C. To encrypt educational content

D. To replace tokenization entirely

Answer = A

Q25. Chunking Challenge

A. Embeddings cannot represent long documents

B. Excessively large chunks dilute semantic specificity

C. Transformers only support images

D. Vector databases reject enterprise files

Answer = B

Q26. Enterprise Recommendation Engine

A. Embedding vectors representing content semantics

B. CPU cache optimization

C. SQL normalization rules

D. Static keyword filtering only

Answer = A

Q27. Hybrid Search Architecture

A. It eliminates all storage requirements

B. It balances exact keyword matching with semantic understanding

C. It disables tokenization overhead

D. It removes transformer dependency completely

Answer = B

Q28. Embedding Drift Scenario

A. Old and new embeddings exist in incompatible vector spaces

B. Transformers stop generating tokens

C. GPU hardware permanently fails during updates

D. Chunking automatically deletes metadata

Answer = A

Q29. Metadata Filtering

A. It reduces semantic retrieval relevance

B. It enables more targeted and context-aware search results

C. It removes embeddings entirely

D. It prevents tokenization

Answer = B

Q30. Enterprise Semantic FAQ System

A. Semantic similarity learned through embeddings

B. Static keyword lookup tables only

C. Spreadsheet formulas

D. Manual hardcoded routing

Answer = A